



1 Record my hand



Ten phone clips, 136 s, of my hand moving an AirPods case into a box; the pinch point tracked in metres.

2 Retarget into LIBERO



Grasp → can, release → basket, the hand's path in between. 500 replays per dataset: A can in place, B moved.

3 Fine-tune SmoIVLA



One policy per dataset, from smolvla_base, which never saw LIBERO; same recipe, 15k steps.

4 Plan with a world model



Policy B samples 8 chunks; the world model imagines 1 s of each and the best-scored one runs.

Experiment 1 can in its usual spot

Test: LIBERO's 50 fixed layouts, can where it always is

Baseline SmoIVLA base, not fine-tuned

0%

Ours A: fine-tuned on my 10 phone clips

90%

Baseline fine-tuned on 44 LIBERO teleop demos

96%

Experiment 2 can moved

Test: can moved 0 to 5 cm, 250 episodes

Ours A: my clips, trained with the can in place

33%

Baseline fine-tuned on 44 LIBERO teleop demos

58%

Ours B: my clips, trained with the can moved

73%

Experiment 3 world model as planner

Test: as in experiment 2, policy B

Ours B alone, no planning

73%

Ours B + world model picks 1 of 8 chunks

85%

Upper bound simulator picks (perfect foresight)

90%